

Parallel FFT: CUDA and Cilk

Ari Bruck


Eric Addison

EEW382V - Parallel Algorithms - Summer 17




Topics

1. Fourier Decomposition
2. Discrete Fourier Transforms
3. Fast Fourier Transforms
 - a. Recursive FFT
 - b. Iterative FFT
 - c. Parallel FFT
4. Parallel FFT using Cilk
5. Parallel FFT using CUDA
6. Performance
7. Conclusion




Fourier Decomposition



Fourier Decomposition



Fourier Decomposition



Fourier Decomposition

$$x(t) = \sum_{k=-\infty}^{\infty} c_k e^{ik\omega_0 t} \longleftrightarrow X[k] = \frac{1}{T} \int_0^T x(t) e^{-ik\omega_0 t} dt$$



$$x(t) = \frac{1}{2\pi} \int_{-\infty}^{\infty} X(\omega) e^{i\omega t} d\omega \longleftrightarrow X(\omega) = \int_{-\infty}^{\infty} x(t) e^{-i\omega t} dt$$

$$x[n] = \frac{1}{2\pi} \int_{-\infty}^{\infty} X(e^{i\omega}) e^{i\omega n} d\omega \longleftrightarrow X(e^{i\omega}) = \sum_{n=-\infty}^{\infty} x[n] e^{-i\omega n}$$

$$x[n] = \frac{1}{N} \sum_{k=0}^{N-1} X[k] e^{i2\pi nk/N} \longleftrightarrow X[k] = \sum_{n=0}^{N-1} x[n] e^{-i2\pi nk/N}$$



Fourier Analysis

$$x_1[n] = \sin(2\pi f_1 n) + 0.5 \cos(2\pi f_2 n) + 0.2 \cos(2\pi f_2 n + \pi/4)$$



Fourier Analysis

$$x_2[n] = U + 0.2 \sin(2\pi f_0 n)$$



Discrete Fourier Transform

$$X[k] = \sum_{n=0}^{N-1} x[n] e^{-i2\pi nk/N}$$

We can implement this on a computer!

DFT Implementation

```
1 DFT(x)
2 N = x.length
3 v = -i*2*pi/N      // i == imaginary unit
4 for k in 1..N
5     X[k] = 0
6     for j in 1.. N
7         X[k] += x[j] * exp(v*k)
```

The Fast Fourier Transform

Q: What is a *Fast* Fourier Transform???

The Fast Fourier Transform

Q: What is a *Fast* Fourier Transform???

A: A Fourier Transform that is not slow!!!

The Fast Fourier Transform

Q: What is a *Fast* Fourier Transform???

A: A Fourier Transform that is not slow!!!

“Fast Fourier Transform” refers to a class of algorithms that can compute a DFT in better than $O(N^2)$ time.

Cooley-Tukey FFT Algorithm

Divide and Conquer!

$$\begin{aligned}X[k] &= \sum_{n=0}^{N-1} x[n]e^{-i2\pi nk/N} \\&= \sum_{n=0}^{N/2-1} x[2n]e^{-i2\pi(2n)k/N} + \sum_{n=0}^{N/2-1} x[2n+1]e^{-i2\pi(2n+1)k/N} \\&= \sum_{n=0}^{N/2-1} x[2n]e^{-i2\pi(2n)k/N} + e^{-i2\pi k/N} \sum_{n=0}^{N/2-1} x[2n+1]e^{-i2\pi(2n)k/N} \\&= \sum_{n=0}^{N/2-1} x_e[n]e^{-i2\pi nk/(N/2)} + e^{-i2\pi k/N} \sum_{n=0}^{N/2-1} x_o[n]e^{-i2\pi nk/(N/2)} \\&= X_e[k] + e^{-i2\pi k/N} X_o[k]\end{aligned}$$

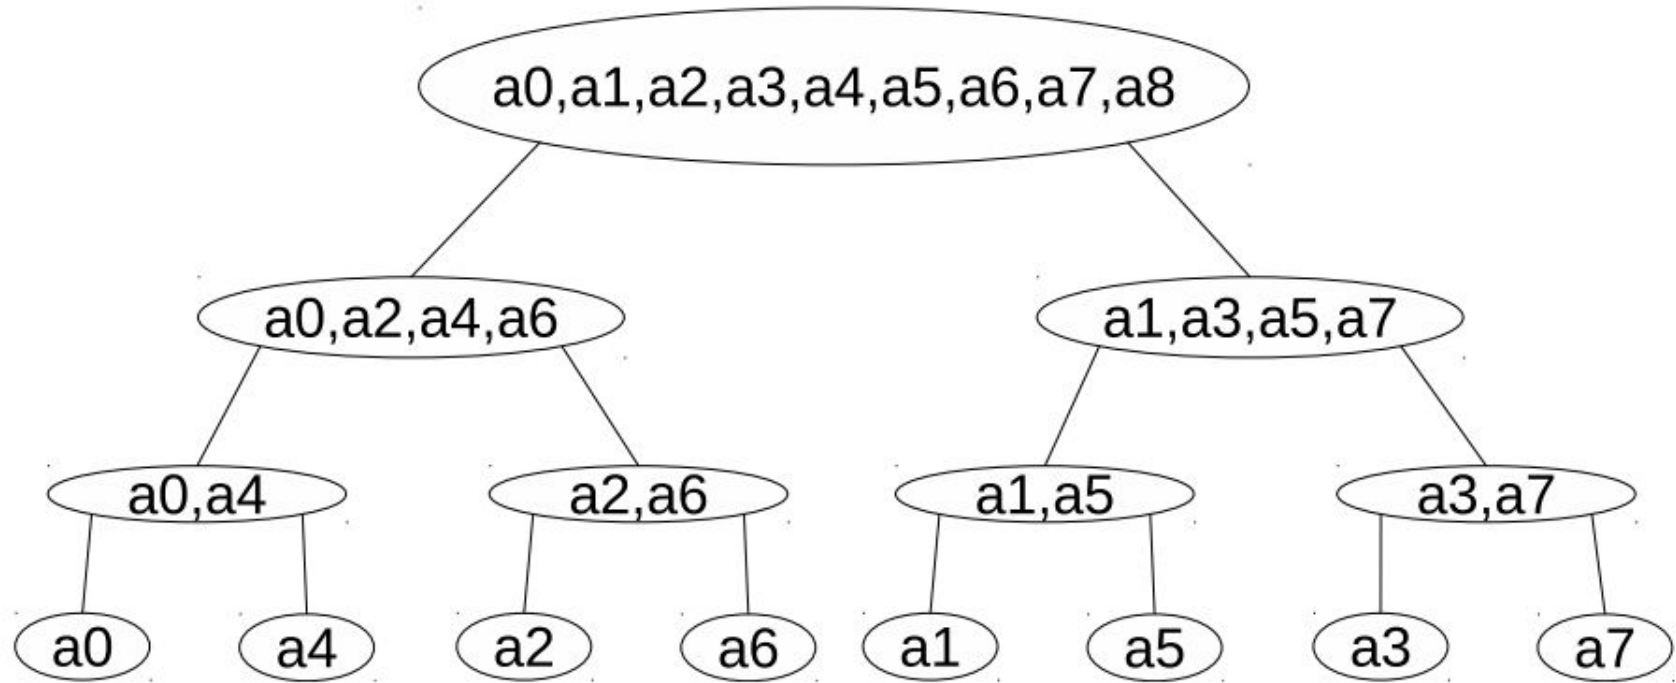
Recursive FFT

```
1  FFT_REC(x)
2      N = x.length
3      if N <= 1
4          return
5      v = -i*2*pi/N      // i == imaginary unit
6      for s in 2..log(N)
7          E = FFT_REC( [x[2*n] for n = 0..N/2-1] )
8          O = FFT_REC( [x[2*n+1] for n = 0..N/2-1] )
9          for j in 0..N/2-1
10             w = exp(v*j)
11             x[2*j] = E[j] + w*O[j]
12             x[2*j+N/2] = E[j] - w*O[j]
13      return x
```

Recursive FFT Time Complexity

$$\begin{aligned}T(N) &= T(N/2) + T(N/2) + \mathcal{O}(N/2) \\&= (T(N/4) + T(N/4) + \mathcal{O}(N/4)) \\&\quad + (T(N/4) + T(N/4) + \mathcal{O}(N/4)) + \mathcal{O}(N/2) \\&= 4T(N/4) + 2\mathcal{O}(N/4) + \mathcal{O}(N/2) \\&= \dots \\&= 2^h T(1) + \sum_{i=1}^h 2^i \mathcal{O}(N/2^i) \\&= N\mathcal{O}(1) + \sum_{i=1}^{\log N} \mathcal{O}(N) \\&= \mathcal{O}(N) + \mathcal{O}(N \log N) \\&= \mathcal{O}(N \log N)\end{aligned}$$

Iterative FFT Algorithm



Iterative FFT: Continued


- Example: 8-element vector = (0,4,2,6,1,5,3,7).
 - Original input (000,001,010,011,100,101,110,111)
 - Bit-Reversed (000,100,010,110,001,101,011,111)

```
1 Bit-Reverse(x)
2     y = empty vector
3     for k in 0..x.length-1
4         k2 = reverse_bits(k)
5         y[k] = x[k2]
6     return y
```

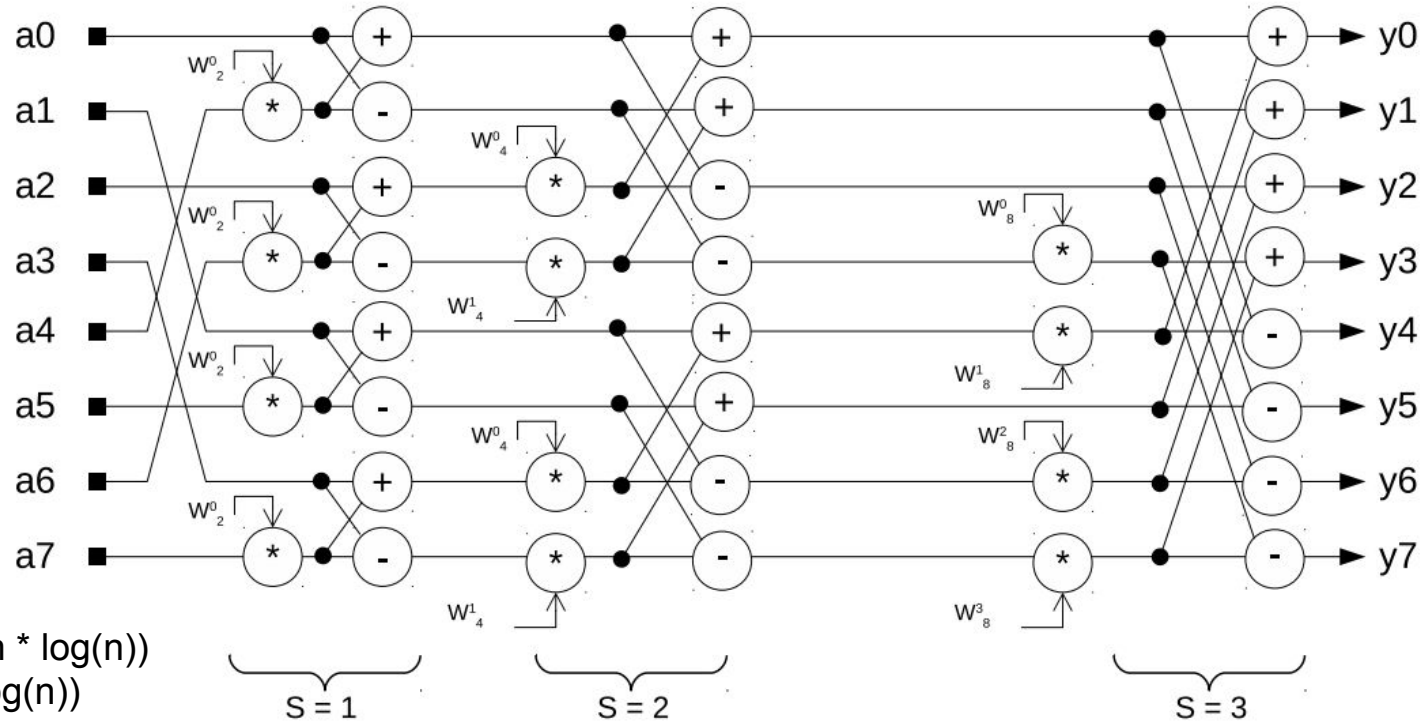
Iterative FFT: Continued

```
1  FFT_ITER(x)
2      Bit-Reverse (x)
3      N = x.length
4      for s in 1..log(N)
5          m = 2^s
6          wm = exp(-i2*pi/m)
7          w = 1
8          for j in 0..(m/2-1)
9              for k in j..(n-1) by m
10                 t = w * x[k + m/2]
11                 u = x[k]
12                 x[k] = u + t
13                 x[k + m/2] = u - t
14             w = w*wm
15      return x
```

FFT Algorithm Performance



Parallel FFT Circuit



$$W(n) = O(n * \log(n))$$

$$T(n) = O(\log(n))$$

Cilk: Introduction

- Cilk is an extension to C/C++ developed by MIT and further enhanced by Intel
- It allows a programmer to easily parallelize a program by inserting certain keywords into targeting portions of an algorithm
- The Cilk framework can easily leverage the hardware capabilities of the system
- The Cilk scheduler is responsible for dividing work between multiple processors
- Parallelizable functions must be exposed by the programmer using a variety of keywords including “spawn” and “sync”
- Cilk enforces rules regarding parent-child relationships. For example, a parent cannot return until all children have “sync[ed]”.

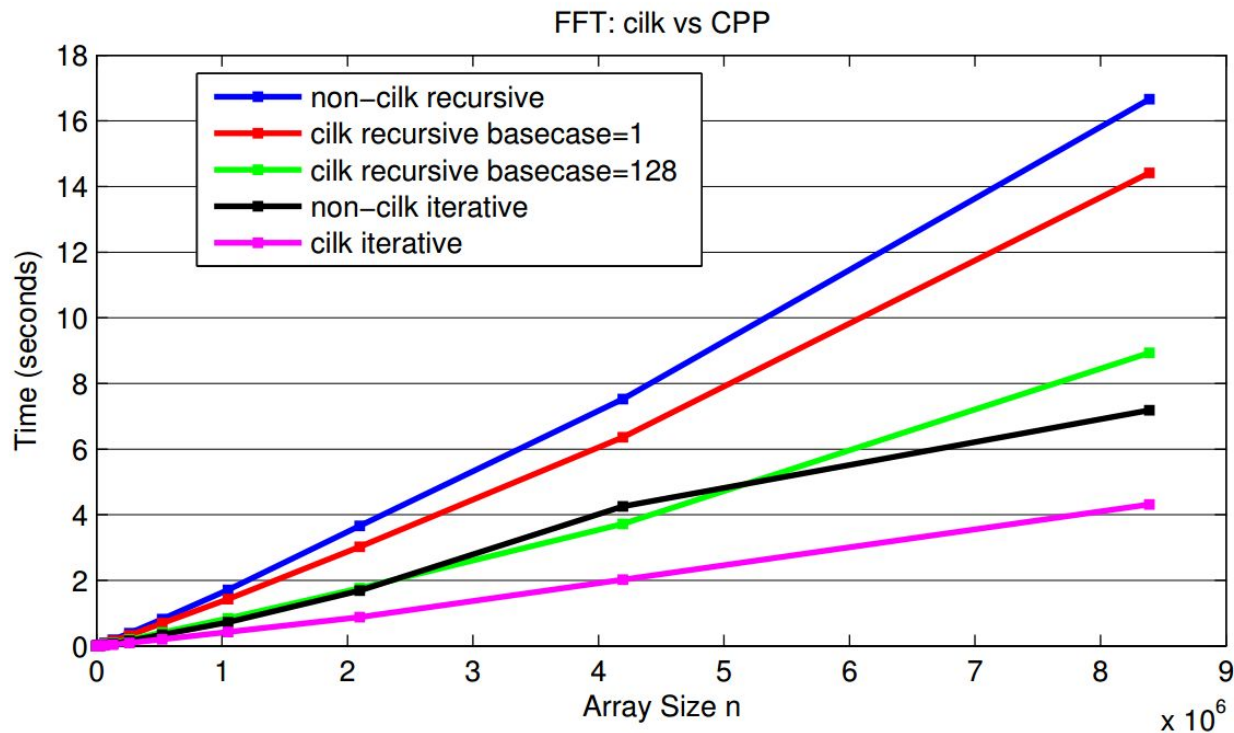
Cilk: Recursive FFT

```
1  FFT_REC(x)
2      N = x.length
3      if N <= 1
4          return
5      v = -i*2*pi/N      // i == imaginary unit
6      for s in 2..log(N)
7          cilk_spawn E = FFT_REC( [x[2*n] for n = 0..N/2-1] )
8          cilk_spawn O = FFT_REC( [x[2*n+1] for n = 0..N/2-1] )
9              cilk_sync
10             cilk_for j in 0..N/2-1
11                 w = exp(v*j)
12                 x[2*j] = E[j] + w*O[j]
13                 x[2*j+N/2] = E[j] - w*O[j]
14     return x
```

Cilk: Iterative FFT

```
1  FFT_ITER(x)
2      Bit-Reverse (x)
3      N = x.length
4      for s in 1..log(N)
5          m = 2^s
6          wm = exp(-i2*pi/m)
7          w = 1
8          cilk_for j in 0..(m/2-1)
9              for k in j..(n-1) by m
10                 t = w * x[k + m/2]
11                 u = x[k]
12                 x[k] = u + t
13                 x[k + m/2] = u - t
14             w = w*wm
15      return x
```


Cilk Performance




Cilk Memory Usage

- Cilk iterative solution uses the same amount of heap space as the C++ solution
 - This might indicate that cilk is using vectorization under the covers with the `cilk_for()` construct
- Cilk recursive solution uses 1.4x more heap space than the C++ solution
 - For test size of 2^{12} it uses ~40MB of heap space vs. ~28MB
- Cilk manipulates the stack, so any stack tracing/profiling induces the “Heisenberg Uncertainty Principle”
 - Act of measuring changes the result

CUDA: Iterative FFT

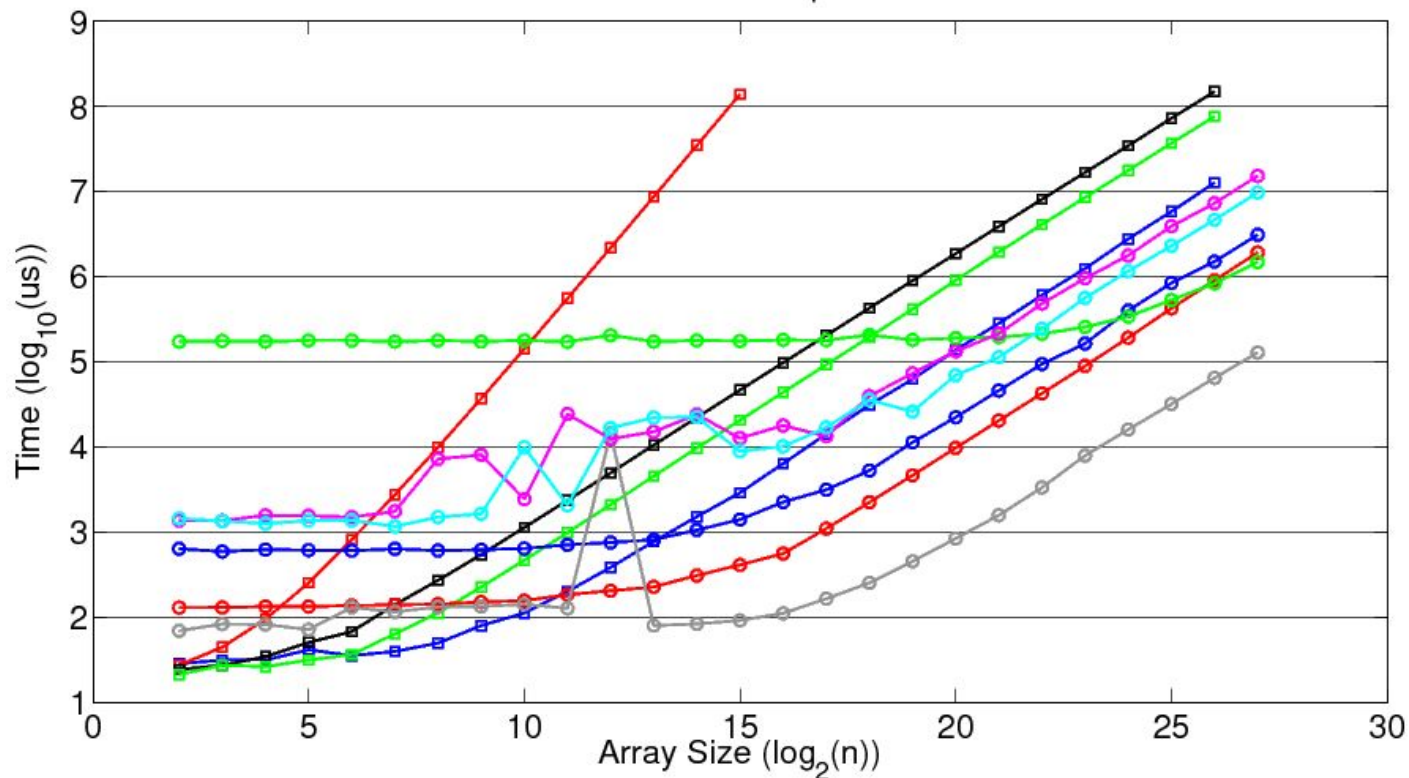
- Load input data from CPU to GPU
- call `bit_reverse_kernel()` to perform array index bit-reversal in global memory
- call `fft_kernel_shared()` to perform the partial FFT using shared memory
- loop over `fft_kernel_finish()` to finish the FFT in global memory
- Retrieve data from GPU memory to CPU

CUDA Performance



Overall Performance

Runtimes for all FFT implementations




Conclusion

- $O(n^2)$ Sucks!
- Recursive Solution is better
 - Modern processors do not optimize well for it
- Iterative Solution is best
- Cilk Parallelism is expensive for small input vectors but scales well
 - Easy in terms of programming
- CUDA is blazing fast, but suffers from memory latency
- CuFFT Algorithm is fast but spends a lot of time “setting up the plan”

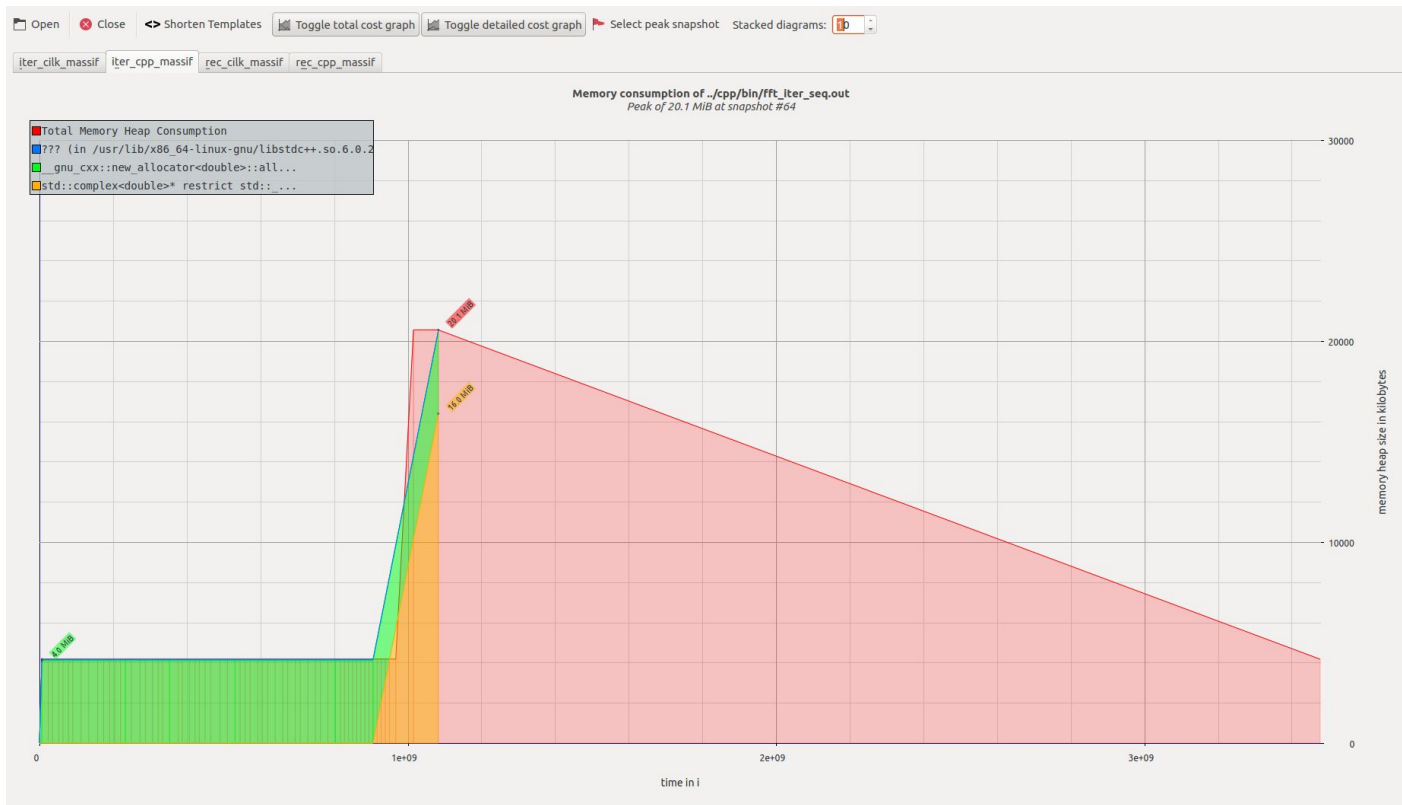
Q&A

Backup Slides


Cilk Iterative FFT Heap Usage



C++ Iterative FFT Heap Usage



Cilk Recursive FFT Heap Usage



C++ Recursive FFT Heap Usage

